


Red-Black Trees

Autumn 2011



RED-BLACK TREES

- *Red-black trees:*

- Binary search trees augmented with node color
- Operations designed to guarantee that the height $h = O(\lg n)$

- First: describe the properties of red-black trees
- Then: prove that these guarantee $h = O(\lg n)$
- Finally: describe operations on red-black trees

RED-BLACK PROPERTIES

- The *red-black properties*:

1. Every node is either red or black
2. Every leaf (NULL pointer) is black
 - Note: this means every “real” node has 2 children
3. If a node is red, both children are black
 - Note: can't have 2 consecutive reds on a path
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES

- Put example on board and verify properties:
 1. Every node is either red or black
 2. Every leaf (NULL pointer) is black
 3. If a node is red, both children are black
 4. Every path from node to descendent leaf contains the same number of black nodes
 5. The root is always black
- *black-height*: # black nodes on path to leaf
 - Label example with h and bh values

HEIGHT OF RED-BLACK TREES

- *What is the minimum black-height of a node with height h ?*
- A: a height- h node has black-height $\geq h/2$
- Theorem: A red-black tree with n internal nodes has height $h \leq 2 \lg(n + 1)$
- *How do you suppose we'll prove this?*

RB TREES: PROVING HEIGHT BOUND

- Prove: n -node RB tree has height $h \leq 2 \lg(n+1)$
- Claim: A subtree rooted at a node x contains at least $2^{bh(x)} - 1$ internal nodes
 - Proof by induction on height h
 - Base step: x has height 0 (i.e., NULL leaf node)
 - *What is $bh(x)$?*

RB TREES: PROVING HEIGHT BOUND

- Prove: n -node RB tree has height $h \leq 2 \lg(n+1)$
- Claim: A subtree rooted at a node x contains at least $2^{bh(x)} - 1$ internal nodes
 - Proof by induction on height h
 - Base step: x has height 0 (i.e., NULL leaf node)
 - *What is $bh(x)$?*
 - A: 0
 - So...subtree contains $2^{bh(x)} - 1$
 $= 2^0 - 1$
 $= 0$ internal nodes (TRUE)

RB TREES: PROVING HEIGHT BOUND

- Inductive proof that subtree at node x contains at least $2^{\text{bh}(x)} - 1$ internal nodes
 - Inductive step: x has positive height and 2 children
 - Each child has black-height of $\text{bh}(x)$ or $\text{bh}(x)-1$ (*Why?*)
 - The height of a child = (height of x) - 1
 - So the subtrees rooted at each child contain at least $2^{\text{bh}(x)-1} - 1$ internal nodes
 - Thus subtree at x contains
$$(2^{\text{bh}(x)-1} - 1) + (2^{\text{bh}(x)-1} - 1) + 1$$
$$= 2 \cdot 2^{\text{bh}(x)-1} - 1 = 2^{\text{bh}(x)} - 1 \text{ nodes}$$

RB TREES: PROVING HEIGHT BOUND

- Thus at the root of the red-black tree:

$$n \geq 2^{\text{bh}(\text{root})} - 1 \quad (\text{Why?})$$

$$n \geq 2^{h/2} - 1 \quad (\text{Why?})$$

$$\lg(n+1) \geq h/2 \quad (\text{Why?})$$

$$h \leq 2 \lg(n+1) \quad (\text{Why?})$$


Thus $h = O(\lg n)$

RB TREES: WORST-CASE TIME

- So we've proved that a red-black tree has $O(\lg n)$ height
- Corollary: These operations take $O(\lg n)$ time:
 - Minimum(), Maximum()
 - Successor(), Predecessor()
 - Search()
- Insert() and Delete():
 - Will also take $O(\lg n)$ time
 - But will need special care since they modify tree

RED-BLACK TREES: AN EXAMPLE

○ *Color this tree:*




Red-black properties:

1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ *Insert 8:*

■ *Where does it go?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ Insert 8

- *Where does it go?*
- *What color should it be?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ Insert 8

- *Where does it go?*
- *What color should it be?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ Insert 11

- *Where does it go?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ Insert 11

- *Where does it go?*
- *What color?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ Insert 11

- *Where does it go?*
- *What color?*
 - Can't be red! (#3)




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ Insert 11

- *Where does it go?*
- *What color?*
 - Can't be red! (#3)
 - Can't be black! (#4)




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ Insert 11

- *Where does it go?*
- *What color?*
- Solution:
recolor the tree




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ Insert 10

- *Where does it go?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION

○ Insert 10

- *Where does it go?*
- *What color?*




1. Every node is either red or black
2. Every leaf (NULL pointer) is black
3. If a node is red, both children are black
4. Every path from node to descendent leaf contains the same number of black nodes
5. The root is always black

RED-BLACK TREES: THE PROBLEM WITH INSERTION


○ Insert 10

- *Where does it go?*
- *What color?*
 - A: no color! Tree is too imbalanced
 - Must change tree structure to allow recoloring
- Goal: restructure tree in $O(\lg n)$ time




RB TREES: ROTATION

- Our basic operation for changing tree structure is called *rotation*:



- *Does rotation preserve inorder key ordering?*
- *What would the code for **rightRotate()** actually do?*


RB TREES: ROTATION



- Answer: A lot of pointer manipulation
 - x keeps its left child
 - y keeps its right child
 - x 's right child becomes y 's left child
 - x 's and y 's parents change
- *What is the running time?*


ROTATION EXAMPLE

- Rotate left about 9:



ROTATION EXAMPLE

- Rotate left about 9:



RED-BLACK TREES: INSERTION

- Insertion: the basic idea
 - Insert x into tree, color x red
 - Only r-b property 3 might be violated (if $p[x]$ red)
 - If so, move violation up tree until a place is found where it can be fixed
 - Total time will be $O(\lg n)$

```
rbInsert(x)
```

```
    treeInsert(x);
```

```
    x->color = RED;
```

```
    // Move violation of #3 up tree, maintaining #4 as invariant:
```

```
    while (x!=root && x->p->color == RED)
```

```
    if (x->p == x->p->p->left)
```

```
        y = x->p->p->right;
```

```
        if (y->color == RED)
```

```
            x->p->color = BLACK;
```

```
            y->color = BLACK;
```

```
            x->p->p->color = RED;
```

```
            x = x->p->p;
```

```
        else    // y->color == BLACK
```

```
            if (x == x->p->right)
```

```
                x = x->p;
```

```
                leftRotate(x);
```

```
            x->p->color = BLACK;
```

```
            x->p->p->color = RED;
```

```
            rightRotate(x->p->p);
```

```
    else    // x->p == x->p->p->right
```

```
        (same as above, but with
```

```
        "right" & "left" exchanged)
```

} Case 1

} Case 2

} Case 3

```
rbInsert(x)
```

```
    treeInsert(x);
```

```
    x->color = RED;
```

```
    // Move violation of #3 up tree, maintaining #4 as invariant:
```

```
    while (x!=root && x->p->color == RED)
```

```
    if (x->p == x->p->p->left)
```

```
        y = x->p->p->right;
```

```
        if (y->color == RED)
```

```
            x->p->color = BLACK;
```

```
            y->color = BLACK;
```

```
            x->p->p->color = RED;
```

```
            x = x->p->p;
```

```
        else // y->color == BLACK
```

```
            if (x == x->p->right)
```

```
                x = x->p;
```

```
                leftRotate(x);
```

```
                x->p->color = BLACK;
```

```
                x->p->p->color = RED;
```

```
                rightRotate(x->p->p);
```

```
    else // x->p == x->p->p->right
```

```
        (same as above, but with
```

```
        "right" & "left" exchanged)
```

} Case 1: uncle is RED

} Case 2

} Case 3

RB INSERT: CASE 1


if (y->color == RED)

 x->p->color = BLACK;


 y->color = BLACK;

 x->p->p->color = RED;

 x = x->p->p;



case 1
.....>



Change colors of some nodes, preserving #4: all downward paths have equal b.h.
The while loop now continues with x's grandparent as the new x

RB INSERT: CASE 1


if (y->color == RED)

 x->p->color = BLACK;

 y->color = BLACK;

 x->p->p->color = RED;


 x = x->p->p;




Same action whether x is a left or a right child

RB INSERT: CASE 2

```
if (x == x->p->right)
    x = x->p;
    leftRotate(x);
// continue with case 3 code
```



case 2
.....>




Transform case 2 into case 3 (x is left child) with a left rotation
This preserves property 4: all downward paths contain same number of black nodes

RB INSERT: CASE 3

x->p->color = BLACK;

x->p->p->color = RED;

rightRotate(x->p->p);



- Case 3:
 - “Uncle” is black
 - Node x is a left child
- Change colors; rotate right

Perform some color changes and do a right rotation

Again, preserves property 4: all downward paths contain same number of black nodes

RB INSERT: CASES 4-6

- Cases 1-3 hold if x 's parent is a left child
- If x 's parent is a right child, cases 4-6 are symmetric (swap left for right)